

Image Processing Basics

In CNN

classmate

Date _____

Page _____

* Image Processing Basics :-

⇒ Image Representation :-

① ⇒ Grayscale Images :→ Understand grayscale images where pixel values represent intensity levels.

→ Grayscale images are represented using a single channel, where pixel values typically range from 0 to 255, representing shades of gray.

→ Each pixel is represented by a single intensity value, where 0 corresponds to black and 255 corresponds to white.

→ Grayscale images are commonly used in image processing tasks where color information is not necessary.

② Color Images :-

→ Color images are represented using multiple channels to capture color information. The most common representation is RGB (Red, Green, Blue).

→ In RGB representation, each pixel is represented by three values, corresponding to the intensity of red, green, and blue channels.

→ Other color representations include HSV (Hue, Saturation, Value), which separates color information into perceptually relevant components.

③ Binary Images :- Binary images are represented using a single bit per pixel, where each pixel is either black (0) or white (1).

→ Binary images are often generated by thresholding grayscale or color images, where pixels above a certain intensity threshold are set to white, and pixels below the threshold are set to black.

→ Binary images are used in tasks such as object detection, segmentation, and feature extraction.

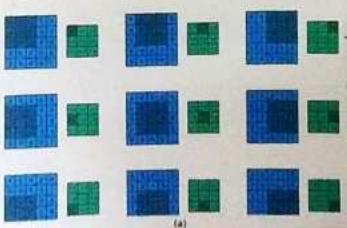
⇒ Image Filtering :-

① Convolution :

→ Convolution is a fundamental operation in image processing where a Kernel (also known as a filter or mask) is applied to an image.

→ The Kernel is a small matrix used to modify the values of pixels in the original image.

→ During Convolution, the Kernel slides over the entire image, and at each position, the sum of element-wise products between the Kernel and the overlapping image region is computed to produce the output pixel value.



② Image Filtering Techniques:-

① Blurring :- → Blurring is a common image filtering technique used to reduce noise and detail in images.

→ It works by averaging pixel values within a local neighborhood defined by the kernel.

→ Popular blurring filters include the Gaussian blur, which assigns higher weights to central pixels and lower weights to surrounding pixels, creating a smoother effect.

⑥ Sharpening :- → Sharpening enhances the edges and details in an image.

→ It works by accentuating the difference between neighbouring pixel values.

→ The sharpening filter typically involves subtracting a blurred version of the image from the original image.

③ Edge Detection:-

- Edge detection is used to identify boundaries within an image where significant changes in intensity occur.
- Sobel, Perwitt, and Canny are common edge detection filters.
- Sobel and Perwitt filters compute the gradient magnitude of the image, highlighting regions of high intensity change.
- The Canny edge detector is a multi-stage algorithm that identifies edges by detecting local maxima of gradient magnitude after applying Gaussian smoothing and gradient calculation.

④ Noise Reduction:-

- Noise reduction filters are used to remove unwanted artifacts or irregularities from images.
- Gaussian filters smooth the image by convolving it with a Gaussian Kernel, which effectively reduces high-frequency noise.
- Median filters replace each pixel's value with the median value in its neighborhood, making them robust to outliers and preserving edges.

⑥ Gaussian Filter :- → The Gaussian Filter is a commonly used image smoothing ^{filter}.

→ It applies a weighted average to the pixels in the image, with the weights determined by a Gaussian function.

The Gaussian filter effectively reduces high-frequency noise in the image while preserving important edges and structures.

→ It is widely used as a preprocessing step before applying other image processing techniques.

The size of the Gaussian Kernel and the standard deviation of the Gaussian function are parameters that affect the smoothing effect.

⑦ Median Filter :- → The median filter replaces each pixel's value with the median value within its neighborhood.

→ It is effective at removing impulsive noise (salt-and-pepper noise) while preserving edges and fine details.

→ Unlike the Gaussian filter, the median filter is non-linear and does not smooth the image in a uniform manner.

It is computationally efficient and robust to outliers in the image.

⑨ Sobel Filter :- → The Sobel filter is used for edge detection in images.

→ It computes the gradient approximation of the image intensity at each pixel.

The filter consists of two separate kernels (one for horizontal changes and one for vertical changes) that highlight edges in these directions.

→ The gradient magnitude at each pixel is computed as the square root of the sum of squares of the horizontal and vertical gradient values.

Sobel filters are widely used due to their simplicity and effectiveness in detecting edges.

(h) Prewitt Filters :- → The Prewitt filter is similar to the Sobel filter and is used for edge detection.

→ It computes the gradient approximation in a similar manner to Sobel but uses different convolution kernels.

Like Sobel, Prewitt filters highlight edges in both horizontal and vertical directions.

→ Prewitt filters are also widely used in image processing applications, especially when Sobel filters may not be appropriate due to specific requirements.

i) Canny Edge Detector :- → The Canny edge detector is a multi-stage algorithm for edge detection.

→ It begins with Gaussian smoothing to reduce noise in the image.

→ The gradient magnitude and direction are then computed using Sobel or Prewitt filters.

Non-maximum suppression is applied to thin the edges to a single-pixel width.

Finally, hysteresis thresholding is used to determine which edges are true edges based on their intensity gradients.

→ The canny edge detector is highly effective in detecting a wide range of edges while minimizing false positives.

⇒ Image Transformations:

① Rotations :- → Rotation involves rotating an image by a certain angle around a specified point.

→ To perform rotation, each pixel's coordinates are transformed based on a rotation matrix.

The rotation matrix is constructed using trigonometric functions (sine & cosine) based on the rotation angle.

→ After applying the rotation matrix, the pixels in the original image are mapped to new coordinates, resulting in a rotated image.

② Scaling :- → Scaling involves resizing an image by a factor along the horizontal and vertical dimensions.

→ To perform scaling, each pixel's coordinates are multiplied by scaling factors along the x and y axes.

→ Scaling can be either uniform (equal scaling along both axes) or non-uniform (different scaling factors along each axis).

③ Translation :-

- Translation involves shifting an image by a certain distance along the horizontal and vertical directions.
- To perform translation, a translation vector specifying the amount of shift along each axis is applied to the coordinates of each pixel.
- Translation does not change the shape or orientation of the image but only shifts its position.

④ Affine Transformations :-

- Affine transformations include combinations of rotation, scaling, translation, and shearing.
- An affine transformation preserves collinearity and parallelism, as well as ratios of distances between points lying on a straight line.
- Affine transformations are represented by a matrix which combines rotation, scaling, and translation matrices into a single transformation matrix.
- Affine transformations allow for more flexibility in modifying the shape, size, and orientation of images compared to individual transformation.

→ Feature extraction :-

① Corner Detection:-

→ Harris Corner Detection →

↳ The Harris corner detector identifies corners in an image by analyzing variations in intensity in different directions.

It computes a corner response function based on the gradient of the image intensity.

↳ Corners are identified as points where the corner response function exceeds a certain threshold.

→ Shi-Tomasi Corner Detector:-

↳ Similar to Harris corner detector, the Shi-Tomasi corner detector also computes a corner response function.

↳ It selects corners based on the minimum eigenvalues of the corner response matrix.

↳ Shi-Tomasi is often preferred over Harris due to its better performance in selecting reliable corners.

② Blob Detection :-

→ Scale-Invariant Feature Transformation (SIFT):

- ↳ SIFT detects Keypoints in an image that are invariant to scale, rotation, and illumination changes.
- ↳ It identifies Keypoints based on local extrema in scale-space.
- ↳ SIFT descriptors are computed for each Keypoint, capturing information about the local image region.

→ Speeded-Up Robust Features (SURF) :-

- ↳ SURF is a fast and efficient alternative to SIFT.
- ↳ It uses integral images to compute image gradients, making it computationally efficient.
- ↳ SURF descriptors are robust to changes in scale and rotation and are used for object recognition and image matching tasks.

→ Image Segmentation :-

- ① Thresholding :- → Thresholding is a simple technique where pixels in an image are classified as belonging to either the foreground or background based on a threshold value.
- It is often used for segmenting objects from the background in grayscale or binary images.
- Common methods include global thresholding, adaptive thresholding, and Otsu's method.

② Region Growing :-

- Region growing starts from seed points and iteratively adds neighboring pixels that satisfy certain criteria (e.g.: similarity in intensity or texture) to form regions.
- It's effective for segmenting regions with uniform properties but may suffer from over-segmentation or under-segmentation based on seed point selection and thresholding criteria.

③ Clustering Algorithm :- (K-means Clustering)

- ↳ K-means partitions the image into 'K' clusters based on pixel intensity or feature similarity.
- ↳ It iteratively assigns pixels to the nearest cluster centroid and updates centroids until convergence.
- ↳ It's widely used but sensitive to initial centroid selection and may not handle non-convex clusters well.

⇒ Image Restoration :-

① Restoration Filters :-

→ Restoration filters aim to remove or reduce noise and enhance image quality.

Common filters include Gaussian filter, median filter, Wiener filter, and bilateral filter.

→ Gaussian filter smooths images by convolving them with a Gaussian Kernel, while removing ~~the~~ noise.

→ Wiener filter and bilateral filter are used for restoring images corrupted by additive noise and preserving edges, respectively.

② Deconvolution :-

- Deconvolution is a process of reversing the effects of convolution, aiming to recover the original image from a degraded one.
- It's used to mitigate blurring caused by motion, defocus, or optical aberrations.
- ~~Regularization~~ deconvolution methods incorporate prior knowledge about the image and noise characteristics to stabilize the inversion process.

③ Blind Deconvolution:-

- Blind deconvolution aims to estimate both the original image and the blur kernel without prior knowledge.
- It's challenging because it requires estimating the unknown blur kernel and the original image simultaneously.
- Techniques include variational methods, sparse priors, and optimization-based approaches.

⇒ Deep learning for Image Processing :-

① Convolutional Neural Networks (CNNs) :-

- CNNs are a class of deep learning models specifically designed for processing structured grid-like data, such as images.
- They consist of convolutional layers, pooling layers, and fully connected layers.
- CNNs are widely used for image classification, object detection, segmentation, and image generation tasks.
- Architecture like AlexNet, VGG, ResNet, and EfficientNet are commonly used for various image processing tasks.

② Applications :-

- Deep learning techniques have been applied to various image processing tasks, ~~included~~ including image classification, object detection, image segmentation, image generation, style transfer, and image super-resolution.

* Convolution and its application in edge detection:-

Edge detection is a fundamental task in image processing aimed at identifying boundaries within images. Convolution is commonly used in edge detection algorithms, with one of the most famous filters, being Sobel filter.

Eg:-

The diagram shows a 6x6 input matrix (n x n) and a 3x3 filter matrix (Filter) [Vertical Edge]. The input matrix is:

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

The filter matrix is:

1	0	-1
1	0	-1
1	0	-1

The output of the convolution step is:

-5	F9
----	----

Calculation for the output value F9:

$$\begin{aligned} & \text{(Like a dot product)} \\ & 3 \times 1 + 0 \times 0 + 1 \times (-1) \\ & + 1 \times 1 + 5 \times 0 + 8 \times (-1) \\ & + 2 \times 1 + 7 \times 0 + 2 \times (-1) \end{aligned}$$

3x3 (f x f)

(Filter) [Vertical Edge]

ON Next Step

Like this we compute the output of the convolution with the desired filter for each step.

→ Important formula for the shape of output matrix ($n_{out} \times n_{out}$):-

$$n_{out} = \underline{n - f + 1}$$

* Padding :-

Padding is a technique used in convolutional neural network (CNN) and other convolution-based operations in image processing and signal processing.

It involves adding extra pixels or values around the boundary of an image or input signal before applying convolution.

→ Purpose of Pad Padding:-

① Preserving Spatial Dimensions: Without padding, the spatial dimensions of the output feature map decrease with each convolutional layer. Padding helps maintain the spatial dimensions of the input and output, which can be crucial for preserving information, especially at the boundary.

② Handling Boundary Effects: When convolving an image, the filter/kernel cannot perfectly cover the edges of the image. Without padding, the information near the edges of the input may be underrepresented in the output feature map. Padding ensures that the convolution operation is applied to all parts of the input image.

$$\text{Padding} = 1$$

on a 3×3 matrix



Making 5×5 resultant matrix.

→ Modified formula for the shape of output matrix ($n_{out} \times n_{out}$):-

$$n_{out} = n + 2p - f + 1$$

⇒ Types of Padding:-

① Valid (No Padding): With valid padding, no padding is added to the input image. As a result, the spatial dimensions of the output feature map are reduced based on the size of the filter and the strides used in the convolution operation.

② Same Padding: Same padding ensures that the output feature map has the same spatial dimensions as the input. Padding is added ~~not~~ evenly to all sides of the input image so that the filter/Kernel fits neatly around the input image.
If the filter size is 'F' and the stride is 's', the amount of padding P can be calculated using the formula:

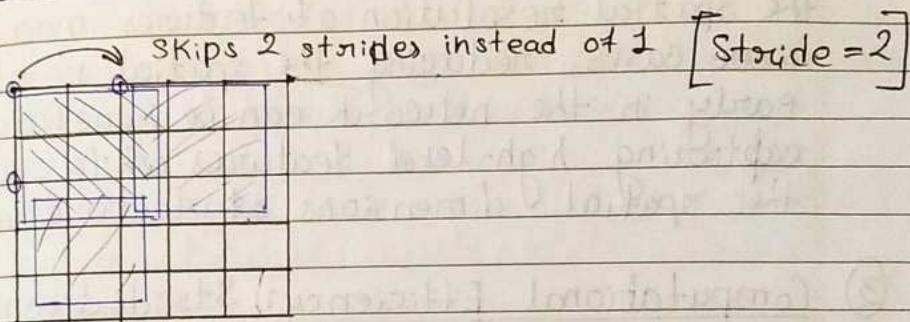
$$P = \left\lfloor \frac{F-1}{2} \right\rfloor \rightarrow (\text{Floor function})$$

* Strided Convolution:-

Strided convolutions are a variation of the standard convolution operation used in deep learning and image processing.

In a regular convolution operation, the filter (or Kernel) moves one pixel at a time across the input image, resulting in an output feature map that preserves spatial dimensions.

Strided convolutions introduce the concept of a stride which dictates the number of pixels the filter moves at each step.



→ Modified formula for the shape of output matrix ($n_{out} \times n_{out}$) :-

$$n_{out} = \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \rightarrow (\text{Floor function})$$

⇒ Purpose and advantages of stride Strided Convolution

- ① Dimensionality Reduction: Strided convolutions allow for downsampling of feature maps. By increasing the stride size, the output feature map's spatial dimensions decrease, providing a way to reduce computational complexity and memory requirements in subsequent layers of the neural network.
- ② Spatial Information Control: Strided convolutions enable control over the spatial resolution of feature maps. In some cases, reducing the spatial resolution early in the network can be beneficial for capturing high-level features while reducing the spatial dimensions of feature maps.
- ③ Computational Efficiency: Strided convolutions can lead to computational efficiency by reducing the number of computations required to process large input images or feature maps.

By skipping over some input pixels based on the stride size, the computational load can be significantly reduced without sacrificing too much information.

* Pooling Layers :-

Pooling layers are commonly used in convolutional neural networks (CNNs) to downsample feature maps, reduce computational complexity, and control overfitting.

→ Pooling Operation: Pooling layers operate on small, contiguous regions of the input feature map, typically referred to as pooling windows or kernels. These windows slide over the input feature map moving by a predefined stride.

→ Pooling Types :-

① Max Pooling :-

In max pooling, the pooling operation selects the maximum value within each pooling window.

This means that the output of each pooling window contains the most prominent feature within the region.

0	2	3	4
1	0	0	0
5	6	0	7
4	1	0	2

2	4
6	7

[2×2 max pooling with stride = 2]

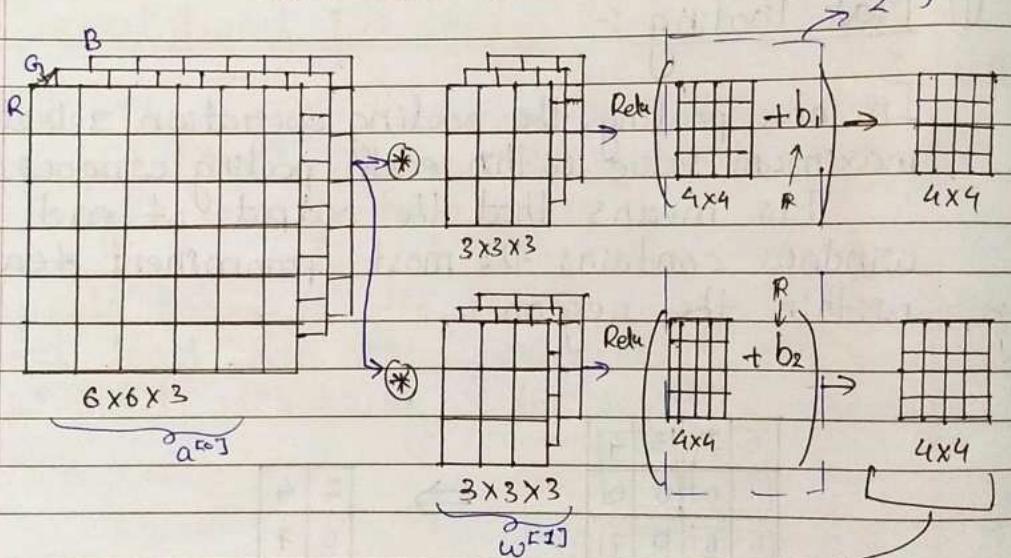
② Average Pooling:-

In average pooling, the pooling operation calculates the average value of all the elements within each pooling window.

This operation smooths out the features and reduces the overall magnitude of the activations.

$\begin{matrix} 1 & 1 & 1 & 0 \\ 0 & 2 & 0 & 1 \end{matrix}$	\rightarrow	$\begin{matrix} 1 & 0.5 \\ 1.5 & 0.5 \end{matrix}$
$\begin{matrix} 2 & 2 & 0 & 0 \\ 1 & 1 & 0 & 2 \end{matrix}$		

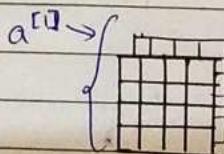
⇒ Example of One layer of Convolutional Network:-



Analogy of z

$$z^{[1]} = w^{[1]} a^{[0]} + b^{[1]}$$

$$a^{[1]} = g(z^{[1]})$$



→ Summary of notation :-

If layer 'l' is a convolution layer:

$f^{[l]}$ = filter size

$p^{[l]}$ = padding

$s^{[l]}$ = stride

$n_c^{[l]}$ = number of filters.

Input : $n_H^{[l-1]} \times n_w^{[l-1]} \times n_c^{[l-1]}$

Output : $n_H^{[l]} \times n_w^{[l]} \times n_c^{[l]}$

$$n_H^{[l]} = \left\lfloor \frac{n_H^{[l-1]} + 2p^{[l]} - f^{[l]} + 1}{s^{[l]}} \right\rfloor; n_w^{[l]} = \left\lfloor \frac{n_w^{[l-1]} + 2p^{[l]} - f^{[l]} + 1}{s^{[l]}} \right\rfloor$$

Each filter: $f^{[l]} \times f^{[l]} \times n_c^{[l-1]}$

Activations: $a^{[l]} \rightarrow n_H^{[l]} \times n_w^{[l]} \times n_c^{[l]}$

Weights : $f^{[l]} \times f^{[l]} \times n_c^{[l-1]} \times n_c^{[l]}$

↑
No. of filters in
layer l

Biases: $n_c^{[l]}$ or $(1, 1, 1, n_c^{[l]})$

⇒ Why Convolutions?

Few reasons:-

① Parameter Sharing :-

A feature detector (such as a vertical edge detector) that's useful in one part of the image is probably useful in another part of the image.

② Sparcity of Connections:-

In each layer, each output value depends only on a small number of inputs.

⇒ Fully Connected layers:-

Fully connected layers are neural network layers in which each neuron is connected to every neuron in the preceding layer, hence the term "fully connected".

These layers are typically found at the end of neural network architectures and are responsible for learning complex patterns and relationships in the input.

Fully connected layers perform classification based on the features extracted by convolutional and pooling layers.

* Image Classification Vs Object Detection Vs Image Segmentation:

Image Classification, Object Detection, and Image Segmentation are three distinct tasks in the field of computer vision, each serving different purposes and requiring different methodologies.

① Image Classification:-

- Image classification involves categorizing an entire image into one of several predefined classes or categories.
- The goal is to assign a single label or class to the entire image based on its content.
- Examples include classifying images of cats, dogs, cars, or handwritten digits.
- Popular datasets for image classification include MNIST, CIFAR-10, and ImageNet.
- Techniques for image classification include traditional machine learning algorithms like SVMs and decision trees, as well as deep learning models like CNNs.

② Object Detection:

- Object detection involves identifying and locating multiple objects within an image along with their corresponding class labels.
- The output typically includes bounding boxes that localize objects and class labels associated with each bounding box.
- Object detection algorithms need to be able to handle multiple objects of varying sizes and orientations within an image.
- Examples include detecting cars, pedestrians, traffic signs, and other objects in images or videos.
- Popular object detection algorithms include R-CNN, Fast R-CNN, Faster R-CNN, SSD, and YOLO

② Image Segmentation :

- Image segmentation involves partitioning an image into multiple segments or regions based on pixel-level classification.
- The goal is to assign a label to each pixel in the image, thereby dividing the image into semantically meaningful regions.
- Image segmentation allows for a more detailed understanding of the spatial structure of objects within an image.
- Examples include segmenting objects from backgrounds, medical image segmentation, and scene parsing.
- Architectures commonly used for image (semantic) segmentation include U-Net, Fully Convolutional Networks (FCNs), and DeepLab.
- Semantic segmentation assigns a single class label to each pixel, while instance segmentation distinguishes between individual object instances within the same class.
- Image Instance Segmentation goes a step further by not only identifying object classes but also distinguishing individual object instances within an image.
Mask R-CNN and its variants are popular architectures for instance segmentation.

⇒ Popular datasets used for image classification tasks include MNIST, CIFAR-10, and ImageNet.

① MNIST: MNIST is a dataset containing 28×28 grayscale images of handwritten digits (0-9). It is widely used as a benchmark dataset for testing various machine learning algorithms, including SVMs, and decision trees.

② CIFAR-10: CIFAR-10 consists of 32×32 color images across 10 classes, such as airplanes, automobiles, birds, cats, etc. It is another popular dataset for image classification tasks, allowing researchers to evaluate the performance of different algorithms.

③ ImageNet: ImageNet is a large-scale dataset containing millions of labeled images across thousands of categories. It has been instrumental in the advancement of deep learning algorithms, particularly Convolutional Network Neural Network (CNNs), for image recognition tasks.

* Classic Networks :-

① 'LeNet' or 'LeNet-5' :-

→ 'LeNet' and 'LeNet-5' refer to the same convolutional neural network architecture.

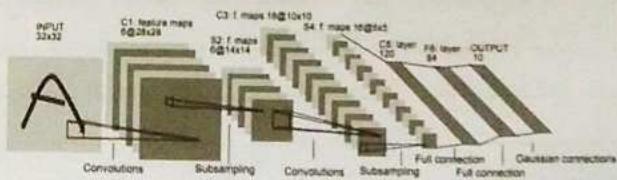
LeNet was one of the pioneering CNN architectures developed by Yann LeCun et al., for handwritten digit recognition, particularly for the task of recognition of the handwritten digits recognition, particularly for the task of recognition in the MNIST dataset.

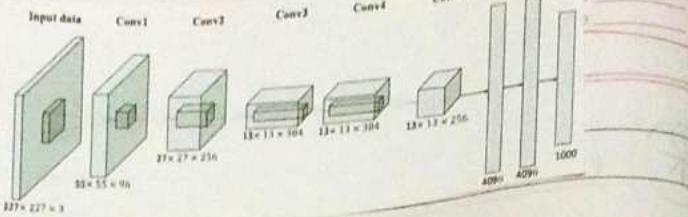
→ The '5' in 'LeNet-5' indicates that the network consists of five layers:

→ It consists of two sets of convolutional and average pooling layers followed by three fully connected layers, which helps in capturing spatial hierarchies and reducing dimensionality.

→ LeNet-5 was designed to be relatively shallow compared to modern networks due to computational constraints at the time.

→ Despite its simplicity, LeNet-5 laid the foundation for subsequent CNN architectures and inspired further research in the field of deep learning.



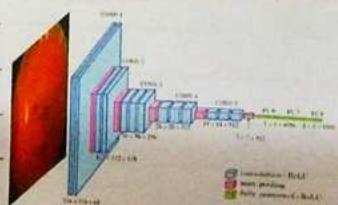


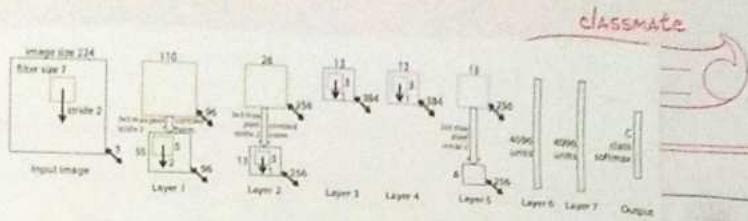
② AlexNet :-

- AlexNet, proposed by Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, won the ImageNet Large Scale Visual Recognition Challenge (ILSVRC) in 2012, significantly advancing the field of Computer Vision.
- It features eight⁽⁸⁾ layers, including five⁽⁵⁾ convolutional layers and three⁽³⁾ fully connected layers, with ReLU activation functions.
- AlexNet employs techniques like data augmentation, dropout, and ReLU activation to improve performance.

③ 'VGGNet' or 'VGGNet - 16' :-

- VGGNet is known for its simplicity and uniformity in architecture. It consists of 16 convolutional layers followed by fully connected layers, with 3x3 filters and max-pooling layers.
- VGGNet is developed by the Visual Geometry Group (VGG) at the University of Oxford, characterized by its deep architecture allowing it to learn complex features but being computationally expensive due to its large number of parameters.





④ ZFNet :-

→ ZFNet, developed by Zeiler and Fergus, won the ILSVRC in 2013, preceding the success of VGGNet and GoogLeNet.

→ It is similar to AlexNet but with modifications such as smaller filter sizes in the first convolutional layer and a smaller stride in the second convolutional layer, which helped improve feature extraction and spatial resolution.

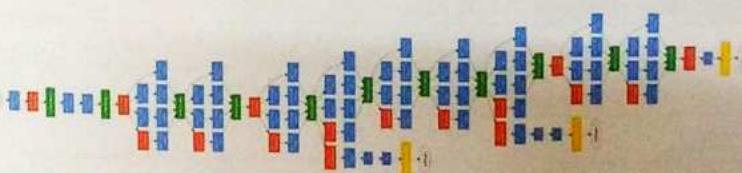
→ ZFNet demonstrated the importance of architectural design choices in achieving superior performance in image classification tasks.

⑤ GoogLeNet :-

→ GoogLeNet, also known as Inception-v1, introduced the concept of inception modules, which allow for efficient use of computational resources by using multiple filter sizes.

→ It consists of 22 layers and is characterized by its deep and wide architecture, achieving high accuracy on the ImageNet dataset.

→ GoogLeNet employs 1×1 convolutions to reduce the dimensionality and computational cost.



* Intermediate CNN Concepts:

① ResNets (Residual Networks):

→ ResNets introduced residual connections, which help mitigate the vanishing gradient problem in very deep neural networks.

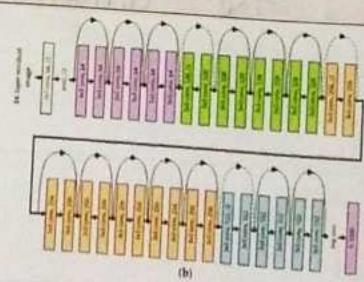
→ Residual connections enable the network to learn residual functions, making it easier to optimize deeper architectures.

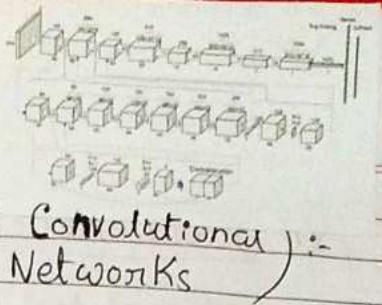
→ Residual connections also known as skip connections. Specifically, a Residual connection involves adding the original input of a layer (or a transformed version of it) to the output of the layer.

This way, the network can learn to adjust the weights of the layer to make its output closer to the desired transformation, but it also has the option to revert to the original input if necessary.

→ Mathematically, if x is the input to a layer, and $F(x)$ represents the transformation applied by the layer, then ~~layer~~ the output of the layer with a residual connection is given by :

$$y = F(x) + x$$



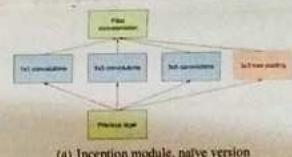


(2) DenseNets (Densely Connected Convolutional Networks) :-

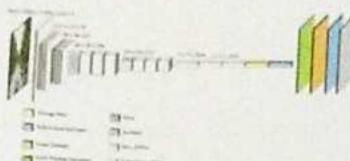
- DenseNets connect each layer to every other layer in a feed-forward fashion within dense blocks.
- DenseNets connectivity encourages feature reuse, improves gradient flow, and reduces the number of parameters.
- DenseNets exhibit strong performance even with relatively shallow architectures and fewer parameters compared to traditional CNNs.

(3) Inception Networks :-

- Inception Networks, popularized by GoogleNet, use inception modules that perform multiple convolutions of different filter sizes and then concatenate the results, to capture spatial hierarchies efficiently.
- The inception modules enable the network to extract features at multiple scales and resolutions, enhancing its representational power.
- Inception Networks strike a balance between computational efficiency and model accuracy, making them suitable for a wide range of computer vision tasks.



(a) Inception module, native version



(4) MobileNet Architecture :-

- MobileNet is designed for mobile and embedded vision applications, optimizing for computational efficiency and low memory footprint.
- It uses depthwise separable convolutions to reduce the number of parameters and computations while maintaining performance.
- MobileNet architectures achieve a good trade-off between model size, accuracy, and inference speed, making them well-suited for deployment on resource-constrained devices.

(5) EfficientNet :-

- EfficientNet introduces a compound scaling method that scales network depth, width, and resolution simultaneously, to achieve optimal performance.
- By scaling all dimensions of the network architecture, EfficientNet achieves state-of-the-art accuracy with significantly fewer parameters and computations compared to traditional CNNs.
- EfficientNet models are highly efficient in terms of both training and inference, making them ideal for practical applications where computational resources are limited.

* Object Localization :-

Object localization is a key task in computer vision that involves identifying the location and extent of objects within an image.

It's a crucial step in various applications such as object detection, facial recognition, and image segmentation.

Key concepts related to object localization:-

① Bounding Boxes :- Bounding boxes are rectangular regions used to represent the location and extent of objects within an image.

→ Representation :- A bounding box is defined by four coordinates:

($x_{\text{min}}, y_{\text{min}}$) for the top-left corner and ($x_{\text{max}}, y_{\text{max}}$) for the bottom-right corner of the box.

→ Purpose : Bounding boxes serve as a simple yet effective way to localize objects in an image, enabling algorithms to identify and differentiate between multiple objects.

② Intersection over Union (IoU) :-

IoU is a metric used to evaluate the accuracy of object localization algorithms.

→ Calculation :- IoU measures the overlap between two bounding boxes by calculating the ratio of the intersection area to the union area of the boxes.

→ Purpose :- IoU provides a quantitative measure of how well a predicted bounding box aligns with the ground truth bounding box. It is commonly used as a criterion for evaluating the performance of object detection and localization models.

③ Landmark Detection :- Landmark detection involves identifying and localizing specific points or landmarks within an object, such as the corners of a person's eyes or the tip of the nose in facial recognition.

→ Application :- Landmark detection is widely used in facial analysis, pose estimation, and medical imaging.

→ Challenges :- Landmark detection algorithms must be robust to variations in pose, scale, and lighting conditions.

② Anchor Boxes :- Anchor boxes, also known as prior boxes, are predefined bounding box shapes used in object detection and localization models.

→ Purpose : Anchor boxes help address the challenge of accurately localizing objects with varying sizes and aspect ratios.

→ Implementation:- In object detection frameworks like YOLO and Faster R-CNN, anchor boxes are used to predict bounding box offsets and dimensions relative to the anchor box shapes.

* Object Detection :-

Object detection is a computer vision task that involves identifying and locating objects within an image or a video.

Several approaches and architectures have been developed to address this task.

Key concepts of Object Detection :

① Sliding Window Detection :- The sliding window detection approach involves systematically moving a window of fixed size across an input image and classifying the content within each window.

→ Process :- A classifier is applied to each window and if the classifier predicts the presence of an object, a bounding box is generated around that region.

→ Limitations :- Computationally expensive as it requires evaluating the classifier for multiple window positions and scales.

② Region-based Approaches : Instead of exhaustively considering all possible windows, region-based approaches focus on generating a set of potential object regions (region proposals) and then classify these regions.

→ Advantages :- More computationally efficient than sliding window approaches as it reduces the number of regions to be evaluated.

→ Example :- Selective Search and EdgeBoxes are algorithms used to generate region proposals.

2.1) R-CNN (Region-based CNN) :- R-CNN integrates the idea of region proposals with deep learning. It uses a region proposal algorithm to propose candidate regions and applies a CNN to each region separately.

→ Training Process :- The model is trained in a two-stage process: pre-training the CNN on a large dataset and fine-tuning it on the target object detection task.

→ Limitations :- Computationally expensive due to separate processing for each region proposal.

2.2) Fast R-CNN :-

→ Improvement :- Fast R-CNN addresses the computational inefficiency of R-CNN by sharing the convolutional features across the entire image, thus avoiding redundant computations for each region proposal.

→ Region of Interest (RoI) Pooling :-

Introduces (RoI) pooling to extract fixed-size features from the shared feature map for each region proposal.

(2.3) Faster R-CNN :- Faster R-CNN further improves efficiency by introducing a Region Proposal Network (RPN) to generate region proposals instead of using external algorithms.

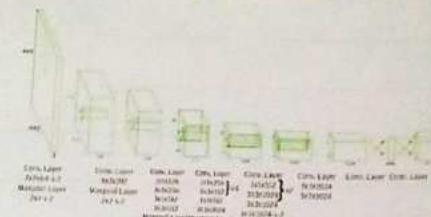
→ End-to-End Training : The entire system is trained end-to-end, making it more streamlined and efficient.

→ Achievements : Faster R-CNN achieves real-time object detection capabilities on a wide range of datasets.

(2) You Only Look Once (YOLO) : YOLO is a single-shot detection model that divides the input image into a grid and predicts boundary bounding boxes and class probabilities directly from the grid cells.

→ Efficiency : YOLO processes the entire image in one forward pass, making it computationally efficient and suitable for real-time application.

→ Anchor Boxes : YOLO uses anchor boxes to predict bounding box offsets and sizes, enabling the detection of objects with various shapes.



④ Non-max Suppression (NMS) :-

→ Concept :- Non-max suppression is a post-processing step that eliminates redundant bounding boxes by keeping only the one with the highest confidence score among overlapping boxes.

→ Implementation :- Used to refine the output of object detection models and avoid multiple detections of the same object.

⑤ Region Proposal Network (RPN) :-

→ The Region Proposal Network (RPN) is a component introduced in the Faster R-CNN architecture.

Unlike traditional region-based approaches that rely on external algorithms to generate region proposals, the RPN is an integral part of the network and learns to generate region proposals directly from the input images.

→ The RPN ~~is~~ is trained only jointly with the object detection network, allowing it to learn to generate highly-quality region proposals based on features extracted from the input image.

⑥ OverFeat in Object Detection:-

- OverFeat was one of the pioneering deep learning models for object detection.
- It uses a sliding window approach combined with convolutional neural networks (CNNs) to classify and localize objects within an image.
- OverFeat extends extracts features at multiple scales using a CNN and then applies a classifier to each sliding window.
- Despite its effectiveness, OverFeat's sliding window approach is computationally intensive.

⑦ Single Shot MultiBox Detector (SSD):-

- SSD is a popular single-stage object detection model known for its simplicity & efficiency.
- Unlike two-stage approaches like Faster R-CNN, SSD directly predicts object bounding boxes and class probabilities from a set of default bounding boxes at multiple scales.
- SSD achieves real-time performance while maintaining competitive accuracy.
- It uses convolutional feature maps of different resolutions to detect objects of varying sizes.

⑧ Mask R-CNN:-

→ Mask R-CNN is a state-of-the-art deep learning model for object detection and instance segmentation, masks etc. It extends the Faster R-CNN architecture by adding a branch for predicting segmentation masks alongside the existing branches for object detection and bounding box regression.

→ It uses a Region Proposal Network (RPN) to generate candidate object proposals based on the features.

The proposals are refined using boundary box regression to obtain more accurate object locations.

→ It adds a parallel branch to the Faster R-CNN architecture, which generates a binary mask for each region proposal.

→ The mask branch operates on the region of interest (RoI) pooled from the feature maps.

→ Mask R-CNN is trained end-to-end using a multi-task loss function.

The loss function includes terms for object classification, bounding box regress (typically smooth L1 loss), and mask prediction (binary cross-entropy)

* Key Topics of Image Segmentation:

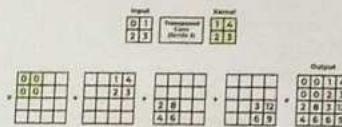
(1) Transpose Convolutions:-

→ Definition: Transpose convolutions, also known as deconvolutions or upsampling convolutions, are used to increase the spatial resolution of feature maps in neural networks.

→ Upsampling: Transpose convolutions perform upsampling by learning to fill in missing information and interpolate between adjacent pixels.

→ Usage: In semantic segmentation tasks, transpose convolutions are commonly used in the decoder network to progressively upsample feature maps and generate high-resolution segmentation masks.

→ Learnable Parameters: Transpose convolutions have learnable parameters, including filter weights and biases, which are optimized during training to improve the quality of upsampling.



② Feature Pyramid Network (FPN):

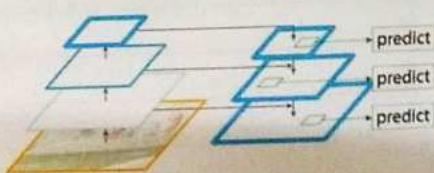
→ Definition: Feature Pyramid Network (FPN) is a multi-scale object detection architecture designed to improve the performance of object detection algorithms, particularly in detecting objects at different scales.

→ Pyramid of Features: FPN generates a pyramid of feature maps by combining feature maps from different layers of a convolutional neural network (CNN).

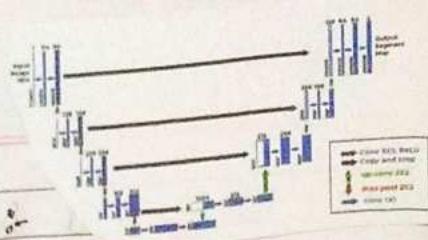
→ Bottom-Up and Top-Down Pathways: FPN utilizes both bottom-up (from lower to higher resolution) and top-down (from higher to lower resolution) pathways to create a feature pyramid.

→ Enhanced Representation: FPN represents enhances the representation of objects at various scales by combining high-level semantic information with detailed spatial information.

→ Usage: FPN is commonly used as a feature extractor in object detection systems, such as Faster R-CNN and RetinaNet, to improve the accuracy and robustness of object localization.



③ U-Net Architecture:



→ Encoder-Decoder Structure: The U-Net architecture comprises an encoder network followed by a decoder network.

→ Purpose: U-Net is specifically designed for biomedical image segmentation tasks, such as cell segmentation and medical image analysis.

→ Contracting Path: The encoder network in U-Net extracts hierarchical features from the input image through a series of convolutional and pooling layers, reducing spatial dimensions.

→ Expansive Path: The decoder network in U-Net uses transpose convolutions to upsample the feature maps and generate pixel-wise predictions.

→ Skip Connections: U-Net incorporates skip connections between corresponding encoder and decoder layers to preserve spatial information and facilitate accurate segmentation.

→ Loss Function: The loss function used in U-Net typically includes terms for pixel-wise classification and regularization.

⇒ Types of Segmentation :-

① Semantic Segmentation :

→ Semantic segmentation assigns a label to every pixel in the image, effectively dividing the image into ~~meaning~~ meaningful regions corresponding to different object classes or categories.

→ Unlike object detection, semantic segmentation does not differentiate between instances of the same class ; it only focuses on classifying each pixel.

② Instance Segmentation :

→ Instance segmentation extends semantic segmentation by not only assigning class labels to pixels but also distinguishing between different instances of the same class.

→ This means that each individual object instance in the image is identified and segmented separately.

③ Panoptic Segmentation:

- Panoptic segmentation combines both semantic and instance segmentation by labeling every pixel with a class label and an instance ID.
- This enables the segmentation of both stuff (e.g., roads, sky) and things (e.g., cars, people) in the image, providing a comprehensive understanding of the scheme.

* Some More Specialized Tasks :

① Behavior Prediction (Multipath):

Behavior prediction, particularly in the context of Artificial Intelligence and Machine Learning, refers to the process of forecasting future actions or behaviors based on historical data and patterns.

The "Multipath" approach likely refers to using multiple pathways or models to predict behavior, which could enhance prediction accuracy and robustness.

MultiPath → Multiple Probabilistic Anchor

Trajectory Hypotheses for Behavior Prediction.

→ Multi-path Approach: Utilizing multiple pathways or models allows for capturing different aspects of behavior, such as short-term and long-term patterns, contextual information, and uncertainty.

These pathways may consist of different types of models, such as recurrent neural networks (RNNs), convolutional neural networks (CNNs), or attention mechanisms.

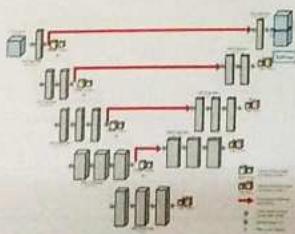
→ Data Representation: Behavior prediction often involves representing input data in a suitable format, such as time-series data, sequential data, or structured data, depending on the nature of the behavior being predicted.

→ Training and Evaluation: Models for behavior prediction are trained using historical data with known outcomes. Evaluation metrics may include accuracy, precision, recall, F1-score, or other domain-specific metrics depending on the application.

② Biomedical Image Segmentation (V-Net):

Biomedical image segmentation is a critical task in medical image analysis that involves partitioning an image into multiple segments or regions of interest.

The "V-Net" likely refers to a specific neural network architecture designed for biomedical image segmentation tasks.



classmate
Date _____
Page _____

→ V-Net Architecture: The V-Net architecture is inspired by the U-Net architecture, which is widely used for biomedical image segmentation.

V-Net typically consists of an encoder-decoder structure with skip connections to facilitate the integration of low-level & high-level features.

→ 3D Convolutional Networks: V-Net is particularly suitable for processing three-dimensional (3D) biomedical images, such as MRI (Magnetic Resonance Imaging) or CT (Computer Tomography) scans. It leverages 3D CNNs to capture spatial information and anatomical structures in the images.

→ Loss Functions: During training, V-Net minimizes a suitable loss function, such as dice loss or cross-entropy loss, to optimize the segmentation accuracy. These loss functions compare the predicted segmentation masks with the ground truth masks.

→ Applications: Biomedical image segmentation with V-Nets has applications in medical diagnosis, treatment planning, image-guided interventions, and medical research.

(3) Hand Detection (Hand Cascader):

Hand detection refers to the process of identifying and localizing human hands within images or video frames. "Hand Cascader" might be a specific algorithm or approach for hand detection, possibly leveraging cascade classifiers or cascade techniques.

→ Cascade Classifiers: Cascade classifiers, popularized by Viola-Jones, are based on the concept of a cascade of classifiers trained to rapidly reject negative image regions that are unlikely to contain the object of interest - in ~~case~~ this case, hands.

→ Feature Selection: Cascade classifiers typically use Haar-like features or other feature descriptions to efficiently capture patterns associated with hands while minimizing computational complexity.

→ Training: Training a hand detector involves providing labeled data (images with annotated hand regions) and optimizing the parameters of the cascade classifier to minimize detection errors and false positives.

→ Real-time Performance: Efficient hand detection algorithms are crucial for real-time applications such as gesture recognition, human-computer interaction, and augmented reality.

④ Face Recognition:-

Face recognition is a computer vision task that involves identifying or verifying the identity of individuals based on their facial features.

It has numerous applications ranging from security and surveillance to access control, personalization, and social media tagging.

→ Face Detection: The first step in face recognition is detecting the presence and location of faces within an image or a video frame.

This involves using algorithms such as Haar cascades, CNNs, or Histogram of Oriented Gradients (HOG) to localize faces.

→ Feature Extraction: Once faces are detected, the next step is to extract meaningful features from the facial images.

Features may include the distances between key facial landmarks, the distribution of pixel intensities, or high-level representations learned by deep neural networks.

→ Deep Learning Approaches: Deep learning methods, particularly CNNs, have significantly advanced the state-of-the-art in face recognition.

CNNs can learn hierarchical representations of facial features directly from raw pixel data, enabling highly accurate face recognition systems.

→ Face Embeddings: Deep learning models used for face recognition often generate embeddings, which are low-dimensional representations of facial features that capture the identity information of individuals.

Face embeddings are typically learned during a training phase and used for comparison during inference.

→ One-shot and Few-shot Learning: In some scenarios, recognition systems need to recognize faces with very few reference examples.

One-shot and few-shot learning techniques enable models to generalize from limited training data and accurately recognize faces with minimal examples.

→ Face Verification Vs. Face Identification:

Face recognition tasks can be divided into face verification (determining whether two faces belong to the same person) and face identification (matching a given face to a database of known faces to determine the person's identity).

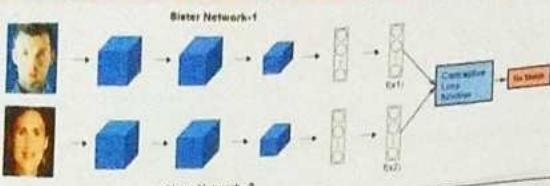
→ Challenges:

Face recognition systems must overcome challenges such as variations in lighting conditions, facial expressions, poses, occlusions, and changes in appearance over time (aging, hairstyle changes, etc.).

→ Privacy and Ethical Considerations:

Face recognition technologies raise concerns related to privacy, surveillance, bias, and ethical use.

Issues such as data privacy, consent, algorithmic fairness, and potential misuse of facial recognition systems are areas of ongoing debate and regulation.



⇒ Siamese Network :-

A Siamese Network is a type of neural network architecture used for tasks involving similarity measurement or verification between pairs of inputs. The basic idea behind a Siamese Network is to use two identical subnetworks, referred to as twin networks or sister networks, which share the same weights and architecture. These networks are typically CNNs in computer vision tasks.

⇒ The Siamese Network architecture is structured as follows:

- ⇒ Two identical shared subnetworks process two separate input samples.
- ⇒ The output embeddings or representations of both inputs are then compared using a distance metric such as Euclidean distance or cosine similarity.
- ⇒ During training, the network learns to minimize the distance between similar pairs of inputs and maximize the distance between dissimilar pairs.

Applications of Siamese Networks include signature verification, face recognition, image similarity comparison, and few-shot learning tasks.

→ Triplet Loss:-

Triplet Loss is a loss function commonly used in Siamese Networks for learning embeddings or representations of input data. It is designed to encourage the network to map similar inputs closer together in the embedding space while pushing dissimilar inputs farther apart.

The Triplet Loss Function is formulated using triplets of samples:

→ Anchor: The input sample for which the network learns to generate a representative embedding

→ Positive: A sample that is similar to the anchor sample

→ Negative: A sample that is dissimilar to the anchor sample.

The objective of Triplet Loss is to minimize the distance between the anchor and positive samples while maximizing the distance between the anchor and negative samples. The loss function can be defined as :

$$\text{Triplet Loss} = \max(d(a,p) - d(a,n) + \alpha, 0)$$

Where :

→ $d(a,p)$ is the distance between the anchor and positive samples

→ $d(a,n)$ is the distance between the anchor and negative samples.

→ α is a margin that defines the minimum difference between the distances of positive & negative pairs.

⇒ Neural Style Transfer:

Neural Style Transfer is a technique that combines the content of one image with the style of another image to generate visually appealing artistic images.

The process involves optimizing a target image to match the style statistics of a style image.

The Convolutional Neural Network (ConvNet) plays a crucial role in Neural Style Transfer by extracting content and style features from the input images. Typically, the content features are extracted from deeper layers of the ConvNet, while the style features are extracted from multiple layers to capture different levels of abstraction.

⇒ The key components of Neural Style Transfer include:

⇒ Content cost function: Measures the difference in content between the target image and the content image.

⇒ Style cost function: Measures the difference in style between the target image and the style image.

⇒ Total variation regularization: Encourages smoothness and reduces noise in the generated image.

→ Content Cost Function:-

The Content Cost Function, also known as the content loss, quantifies the difference in content between the target image G_t and the content image C .

It is calculated based on the feature maps extracted from a pre-trained ConvNet, such as VGGNet, using a specific layer's activations.

The content cost function $J_{\text{Content}}(C, G_t)$ is typically computed as the mean squared error (MSE) between the feature representations of the content image $\text{conv}_l(C)$ and the target image $\text{conv}_l(G_t)$ at a chosen layer l of the ConvNet:

$$J_{\text{Content}}(C, G_t) = \frac{1}{2} \|\text{conv}_l(C) - \text{conv}_l(G_t)\|_2^2$$

where:

- $\text{conv}_l(C)$ and $\text{conv}_l(G_t)$ are the feature representations of the content and target images at layer l , respectively.



⇒ Style Cost Function:

The Style Cost Function, also known as the style loss, quantifies the difference in style between the target image G and the style image S .

It is computed based on the statistics of feature maps extracted from multiple layers of a pre-trained ConvNet.

The style cost function $J_{\text{style}}(S, G)$ is calculated using the Gram matrices of the feature representations, which capture the correlations between different feature maps:

$$J_{\text{style}}(S, G) = \sum_l \frac{\lambda_l}{4 \times n_l^2 \times m_l^2} \| \text{Gram}_l(S) - \text{Gram}_l(G) \|^2$$

Where :

- $\text{Gram}_l(S)$ and Gram_l are the Gram matrices of the style image ' S ' and the target image ' G ' at layer ' l ', respectively.
- ' n_l ' is the number of feature maps, and ' m_l ' is the height times width of each target feature map at layer l .
- ' λ_l ' is a weighting factor for the contribution of each layer to the total style loss.

→ 1D and 3D Generalizations:

1D and 3D generalizations refer to extending certain neural network architectures or operations from 2D (images) to 3D sequences or 3D (Volumes).

→ 1D Convolutions: Instead of convolving over 2D spatial dimensions (width and height), 1D convolutions are applied along the temporal axis or sequence dimension.

They are commonly used in tasks such as time series analysis, natural language processing (NLP), and audio processing.

→ 3D Convolutions: 3D convolutions extend the concept of 2D convolutions to three-dimensional volumes, incorporating depth as an additional dimension.

They are widely used in video analysis, medical imaging and volumetric data processing tasks.

These generalizations allow neural networks to process and extract features from different types of data, facilitating the development of models for diverse applications ~~etc~~ beyond traditional image processing.

→ Some More important concepts:-

→ Object Tracking: Object tracking refers to the process of locating and following objects in a video sequence over time. It involves identifying and monitoring the movement of objects as they move through a scene. Object tracking has numerous applications, including surveillance, video analysis, autonomous vehicles, human-computer interaction, and augmented reality.

→ Localization: Localization is the process of determining the position of an object or entity within a given space.

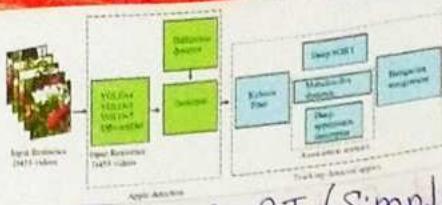
In the context of object tracking, localization involves pinpointing the exact location of an object within each frame of a video sequence. This can be achieved using various techniques such as feature extraction, template matching, and deep learning-based methods.

→ Motion: Motion refers to the movement or change in position of objects within a scene over time. In object tracking, motion analysis is a fundamental aspect used to track the movement of objects from one frame to another. Understanding motion patterns helps in predicting future positions of objects and improving the accuracy of object tracking algorithms.

→ Flow of Optics: Optical flow refers to the pattern of apparent motion of objects within a visual scene caused by the relative motion between the observer (camera) and the objects. Optical flow algorithms estimate the motion of objects by analyzing the pixel intensity changes between consecutive frames in a video sequence. Optical flow is often used as a cue for object tracking and motion analysis.

→ Motion Vector: A motion vector represents the direction and magnitude of motion between two consecutive frames in a video sequence. Motion vectors are commonly used in video compression and object tracking algorithms to describe the motion of objects within a scene. They provide valuable information for predicting the movement of objects and improving the efficiency of object tracking systems.

→ Tracking & Features: Tracking features are distinctive characteristics or attributes of objects that are used to track them across multiple frames in a video sequence. These features can include color, texture, shape, edges, corners, and Keypoints. Feature-based tracking algorithms extract and match these features between frames to identify and track objects accurately.



→ Deep SORT : Deep SORT (Simple Online and Realtime Tracking with a Deep Association Metric) is a state-of-the-art object tracking algorithm that integrates deep learning-based object detection with a sophisticated tracking framework.

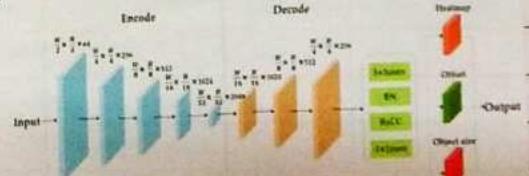
It combines a deep neural network for object detection (such as YOLO) with a SORT algorithm for online multi-object tracking. Deep SORT achieves high accuracy and real-time performance in tracking multiple objects simultaneously.

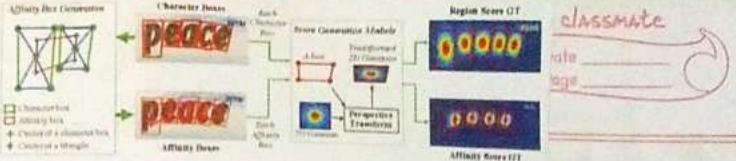
* Specialized Architectures :-

① CenterNet :- CenterNet is a state-of-the-art object detection framework designed to detect objects as keypoints. It directly predicts object centers and regresses bounding boxes around them.

→ Unlike traditional object detection methods that use anchor boxes or region proposal networks, CenterNet directly predicts objects' centers and sizes.

→ CenterNet achieves high accuracy while maintaining efficiency, making it suitable for real-time applications.





② CRAFT: → CRAFT (Character Region Awareness for Text Detection) is a text detection model designed to accurately detect text regions in natural images.

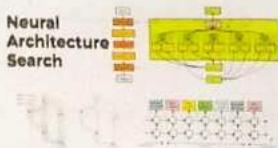
→ It leverages deep learning techniques to localize and segment text regions with high precision.

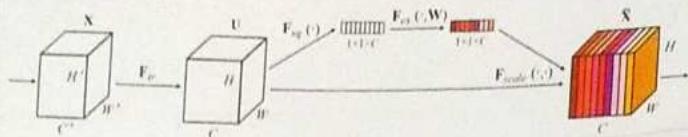
→ CRAFT is particularly effective in handling irregular and curved text, making it suitable for various text detection tasks in real-world scenarios.

③ NAS (Neural Architecture Search): → NAS is a technique that automates the design of neural network architectures, including CNNs, by searching through a predefined space of network architectures.

→ It typically involves using reinforcement learning, evolutionary algorithms, or gradient-based optimization methods to search for architectures that optimize ~~a~~ a specific objective, such as accuracy or efficiency.

→ NAS has led to the discovery of novel architectures that outperform handcrafted designs in various tasks, including image classification, object detection, and semantic segmentation.





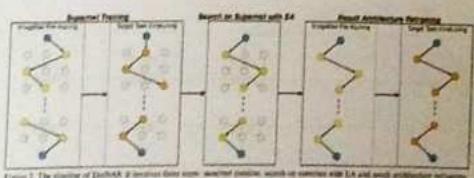
④ SENet (Squeeze-and-Excitation Networks):

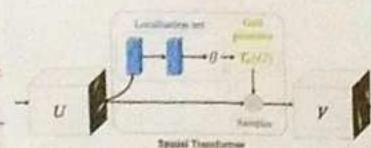
- SENet introduces a mechanism called "channel-wise attention" to adaptively recalibrate channel-wise feature responses.
- It learns to emphasize informative features and suppress irrelevant ones by modeling interdependencies between channels.
- SENet enhances the representational power of CNNs and improves their performance across various computer vision tasks, including image classification and object detection.

⑤ DetNas : → DetNas is a recent advancement in object detection that combines Neural Architecture Search (NAS) with object detection frameworks.

→ It automatically searches for optimal architectures tailored to the object detection task, leading to improved accuracy and efficiency.

→ DetNas explores different design choices, including network depth, layer configurations, and feature fusion strategies, to optimize the overall performance of object detection models.





* Attention Mechanisms:

① Spatial Transformer Networks (STN): → STNs are neural network components that learn to perform spatial transformations on input images in a differentiable manner.

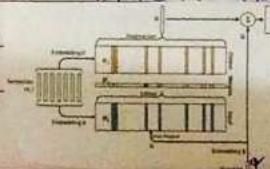
- They can dynamically transform images by learning to focus on relevant regions and ignore irrelevant ones, improving the model's robustness to variations in scale, rotation, and translation.
- STNs have applications in various tasks such as image classification, object localization, and geometric data augmentation.

② End - to - End Memory Networks: → End - to - End Memory

Networks are neural architectures equipped with an external memory component that allows them to store and retrieve information over long sequences.

- They excel in tasks that require reasoning over multiple steps or handling long-term dependencies, such as question answering and language understanding.

→ End - to - End Memory Networks leverage attention mechanisms to select selectively read from and write to memory, enabling them to perform complex inference and reasoning tasks.



(2) Grad-CAM (Gradient-weighted Class Activation Mapping):

→ Grad-CAM is a technique for visualizing and interpreting CNNs by generating heatmaps that highlight regions of input images responsible for specific predictions.

→ It computes the gradient of the target class's score with respect to feature maps in the final convolutional layer, which indicates the importance of each feature map for the prediction.

→ Grad-CAM provides insights into CNN decision-making processes and helps identify which parts of the input image contribute most to the model's predictions.

(4) Highway Networks: → Highway Networks mitigate the vanishing gradient problem and enable the training of deeper networks with improved performance.

→ Highway Networks are neural architectures designed to facilitate the training of very deep networks by introducing "highways" that enable information to flow more easily through the network.

→ They incorporate gating mechanisms inspired by Long Short-Term Memory (LSTM) networks, allowing the model to selectively carry information across layers.

