



Style transfer



Source image (**Style**)




Target image (**Content**)



Output ([deepart](#))

A Neural Algorithm of Artistic Style [[Gatys et al. 2015](#)]

Medical imaging



3D imaging
MRI, CT



Image guided surgery
[Grimson et al., MIT](#)

Spatial and Intensity Resolution

- Spatial resolution

- A measure of the smallest discernible detail in an image
- stated with *line pairs per unit distance*, *dots (pixels) per unit distance*, *dots per inch (dpi)*


- Intensity resolution

- The smallest discernible change in intensity level
- stated with *8 bits*, *12 bits*, *16 bits*, *etc.*

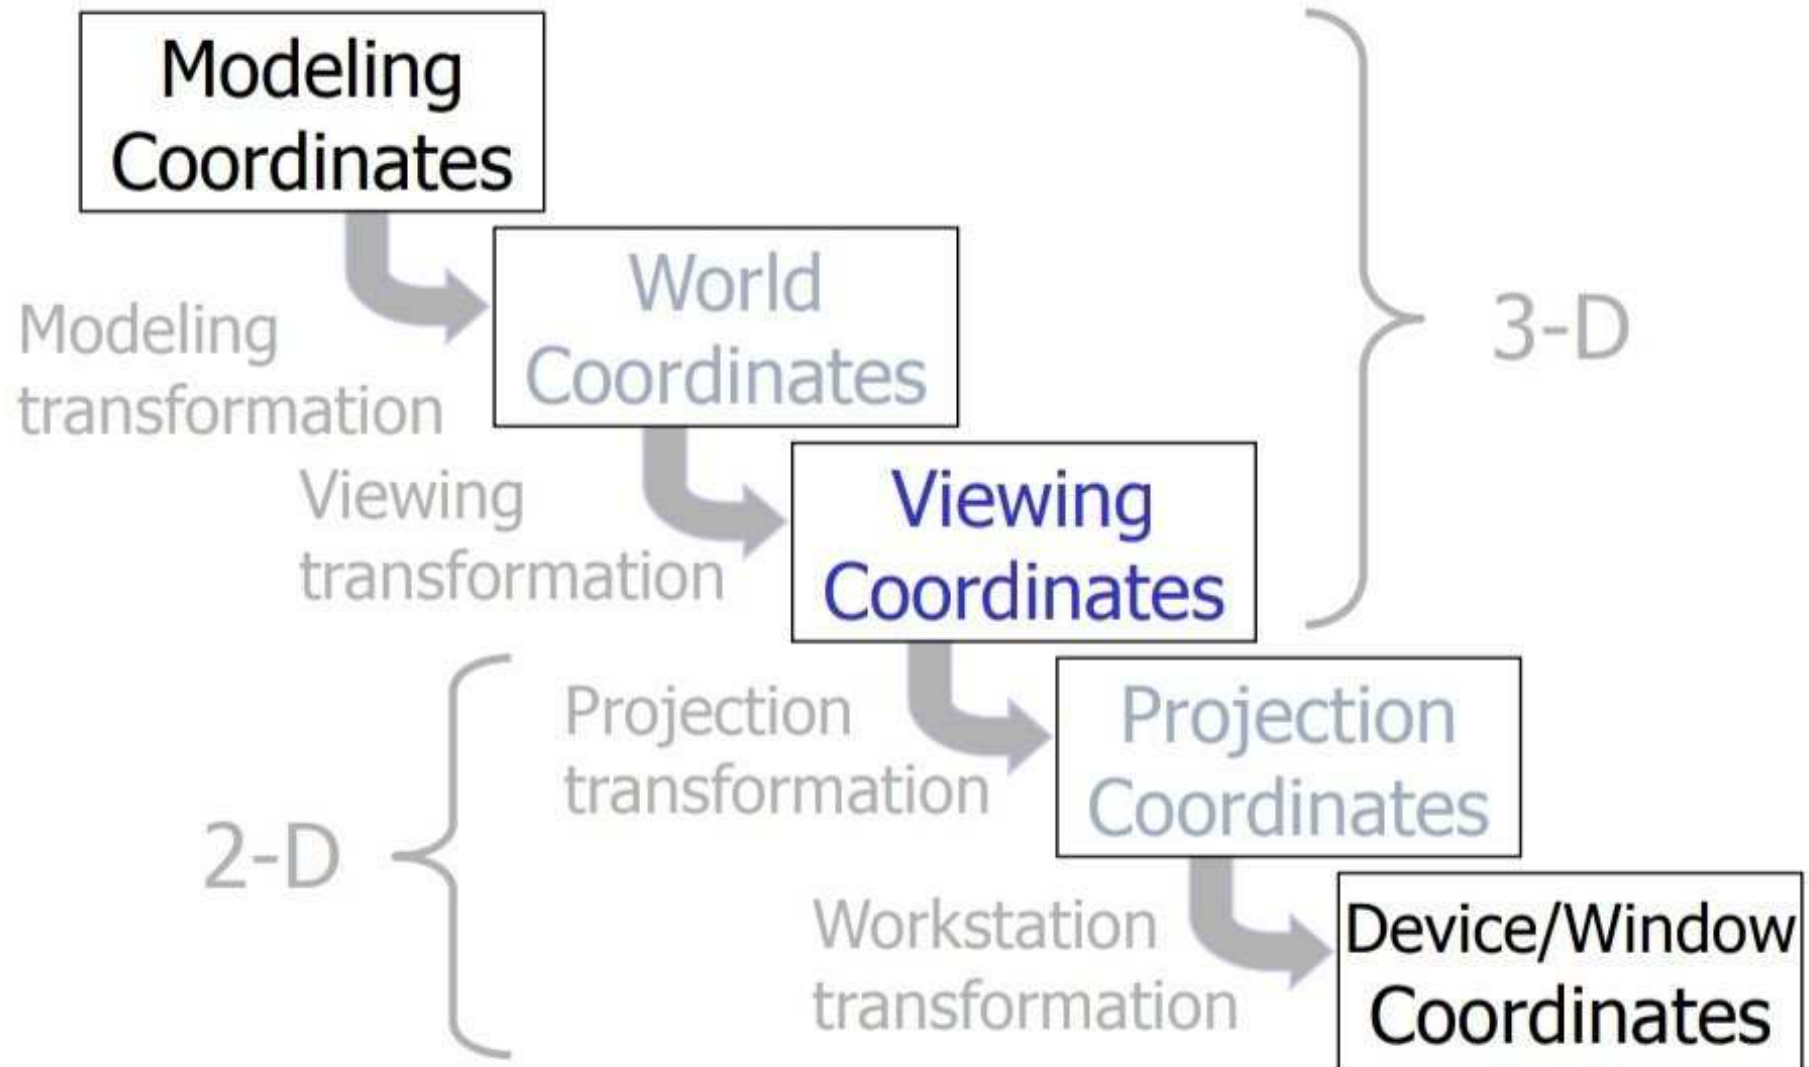
Difference ?

The Eye

- The human eye is a camera!
 - **Iris** - colored annulus with radial muscles
 - **Pupil** - the hole (aperture) whose size is controlled by the iris
 - What's the "film"? photoreceptor cells (rods and cones) in the **retina**




3-D Viewing Process




Projection

- General definition
 - Transform points in n -space to m -space ($m < n$)
- In computer graphics
 - Map viewing coordinates to 2D screen coordinates




Perspective Projection

- Compute 2D coordinates from 3D coordinates with similar triangles



Perspective Projection

- Compute 2D coordinates from 3D coordinates with similar triangles



Question 5

- In the following arrangement of pixels, what's the value of the chessboard distance between the circled two points?

0	0	0	0	0
0	0	1	1	0
0	1	1	0	0
0	1	0	0	0
0	0	0	0	0
0	0	0	0	0

Question 6

- In the following arrangement of pixels, what's the value of the city-block distance between the circled two points?

0	0	0	0	0
0	0	1	1	0
0	1	1	0	0
0	1	0	0	0
0	0	0	0	0
0	0	0	0	0

Mathematical Operations in DIP

- Linear vs. Nonlinear Operation

$$H[f(x, y)] = g(x, y)$$

$$H[a_i f_i(x, y) + a_j f_j(x, y)]$$

$$= H[a_i f_i(x, y)] + H[a_j f_j(x, y)]$$

$$= a_i H[f_i(x, y)] + a_j H[f_j(x, y)]$$

$$= a_i g_i(x, y) + a_j g_j(x, y)$$



Additivity

Homogeneity


H is said to be a **linear operator**;

H is said to be a **nonlinear operator** if it does not meet the above qualification.

Spatial Domain Process

Intensity transformation function

$$s = T(r)$$



a b

FIGURE 3.2

Intensity transformation functions.

(a) Contrast-stretching function.

(b) Thresholding function.

Image Negatives





Image negatives


$$s = L - 1 - r$$

Log Transformations



Log Transformations
 $s = c \log(1 + r)$

Example: Log Transformations




a b

FIGURE 3.5
(a) Fourier spectrum.
(b) Result of applying the log transformation in Eq. (3.2-2) with $c = 1$.

use of log transform?
and other transforms?


Power-Law (Gamma) Transformations



$$s = cr^\gamma$$

FIGURE 3.6 Plots of the equation $s = cr^\gamma$ for various values of γ ($c = 1$ in all cases). All curves were scaled to fit in the range shown.

Example: Gamma Transformations



Cathode ray tube (CRT) devices have an intensity-to-voltage response that is a power function, with exponents varying from approximately 1.8 to 2.5

$$S = r^{1/2.5}$$

Piecewise-Linear Transformations

- **Contrast Stretching**
 - Expands the range of intensity levels in an image so that it spans the full intensity range of the recording medium or display device.
- **Intensity-level Slicing**
 - Highlighting a specific range of intensities in an image often is of interest.

Bit-plane Slicing




FIGURE 3.13
Bit-plane
representation of
an 8-bit image.

Bit-plane Slicing



a	b	c
d	e	f
g	h	i

FIGURE 3.14 (a) An 8-bit gray-scale image of size 500×1192 pixels. (b) through (i) Bit planes 1 through 8, with bit plane 1 corresponding to the least significant bit. Each bit plane is a binary image.

Histogram Processing

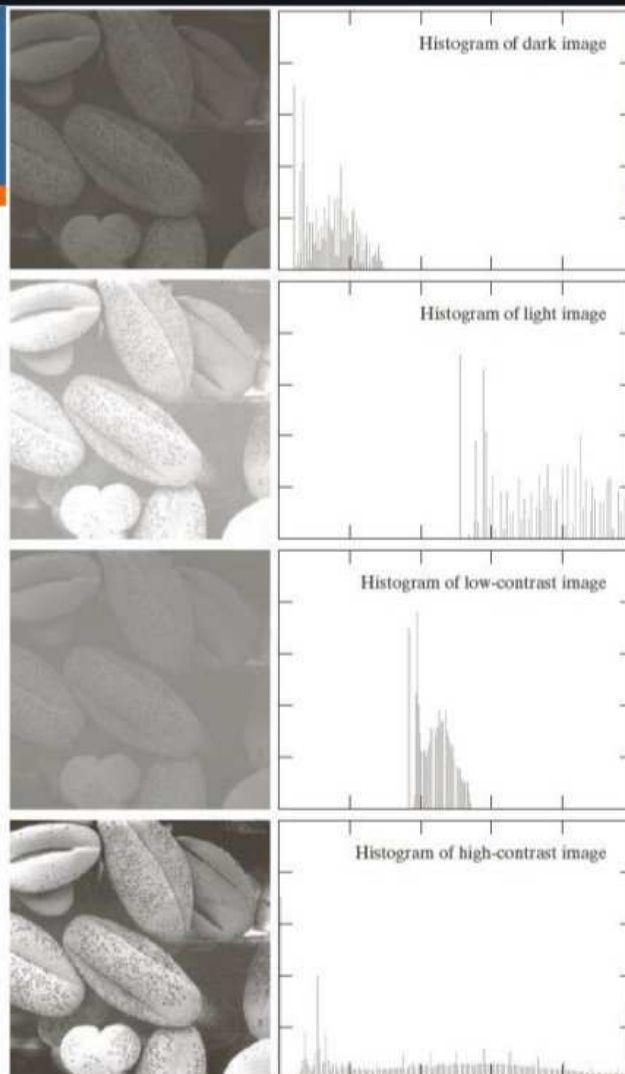
Histogram $h(r_k) = n_k$

r_k is the k^{th} intensity value

n_k is the number of pixels in the image with intensity r_k

Normalized histogram $p(r_k) = \frac{n_k}{MN}$

n_k : the number of pixels in the image of size $M \times N$ with intensity r_k



Histogram Equalization

The intensity levels in an image may be viewed as random variables in the interval $[0, L-1]$.

Let $p_r(r)$ and $p_s(s)$ denote the probability density function (PDF) of random variables r and s .




FIGURE 3.18 (a) An arbitrary PDF. (b) Result of applying the transformation in Eq. (3.3-4) to all intensity levels, r . The resulting intensities, s , have a uniform PDF, independently of the form of the PDF of the r 's.




FIGURE 3.20 Left column: images from Fig. 3.16. Center column: corresponding histogram-equalized images. Right column: histograms of the images in the center column.

*Benefit of
histogram
equalization?*

Order-statistic (Nonlinear) Filters

- Nonlinear
- Based on ordering (ranking) the pixels contained in the filter mask
- Replacing the value of the center pixel with the value determined by the ranking result

E.g., median filter, max filter, min filter

Use ? Effects ?

Image Enhancement in Frequency Domain

*Little
overview*

Fundamentals

- Let R represent the entire spatial region occupied by an image. Image segmentation is a process that partitions R into n sub-regions, R_1, R_2, \dots, R_n , such that

(a) $\bigcup_{i=1}^n R_i = R.$

(b) R_i is a connected set. $i = 1, 2, \dots, n.$

(c) $R_i \cap R_j = \Phi.$

(d) $Q(R_i) = \text{TRUE}$ for $i = 1, 2, \dots, n.$

(e) $Q(R_i \cup R_j) = \text{FALSE}$ for any adjacent regions R_i and $R_j.$

Definition of image segmentation


Detection of Isolated Points

- The Laplacian

$$\nabla^2 f(x, y) = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

$$= f(x+1, y) + f(x-1, y) + f(x, y+1) + f(x, y-1) - 4f(x, y)$$

$$g(x, y) = \begin{cases} 1 & \text{if } |R(x, y)| \geq T \\ 0 & \text{otherwise} \end{cases} \quad R = \sum_{k=1}^9 w_k z_k$$



a
b c d

FIGURE 10.4

(a) Point detection (Laplacian) mask. (b) X-ray image of turbine blade with a porosity. The porosity contains a single black pixel. (c) Result of convolving the mask with the image. (d) Result of using Eq. (10.2-8) showing a single point (the point was enlarged to make it easier to see). (Original image courtesy of X-TEK Systems, Ltd.)

Detecting Line in Specified Directions

-1	-1	-1	2	-1	-1	-1	2	-1	-1	-1	2
2	2	2	-1	2	-1	-1	2	-1	-1	2	-1
-1	-1	-1	-1	-1	2	-1	2	-1	2	-1	-1
Horizontal			+45°			Vertical			-45°		

FIGURE 10.6 Line detection masks. Angles are with respect to the axis system in Fig. 2.18(b).

- Let R_1 , R_2 , R_3 , and R_4 denote the responses of the masks in Fig. 10.6. If, at a given point in the image, $|R_k| > |R_j|$, for all $j \neq k$, that point is said to be more likely associated with a line in the direction of mask k .

Edge Detection

- Edges are pixels where the brightness function changes abruptly
- Edge models




FIGURE 10.8

From left to right, models (ideal representations) of a step, a ramp, and a roof edge, and their corresponding intensity profiles.

z_1	z_2	z_3
z_4	z_5	z_6
z_7	z_8	z_9

-1	0	0	-1
0	1	1	0

Roberts

-1	-1	-1	-1	0	1
0	0	0	-1	0	1
1	1	1	-1	0	1

Prewitt

-1	-2	-1	-1	0	1
0	0	0	-2	0	2
1	2	1	-1	0	1

Sobel

a
b c
d e
f g

FIGURE 10.14
A 3×3 region of an image (the z 's are intensity values) and various masks used to compute the gradient at the point labeled z_5 .

0	1	1	-1	-1	0
-1	0	1	-1	0	1
-1	-1	0	0	1	1

Prewitt

0	1	2	-2	-1	0
-1	0	1	-1	0	1
-2	-1	0	0	1	2

Sobel

a b
c d

FIGURE 10.15
Prewitt and Sobel masks for detecting diagonal edges.

Image Classification pipeline



An image classifier

```
def classify_image(image):  
    # Some magic here?  
    return class_label
```

Unlike e.g. sorting a list of numbers,

no obvious way to hard-code the algorithm for recognizing a cat, or other classes.


```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Nearest Neighbor classifier

Memorize training data

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Nearest Neighbor classifier

For each test image:
Find closest train image
Predict label of nearest image

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Nearest Neighbor classifier

Q: With N examples,
how fast are training
and prediction?

```
import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred
```

Nearest Neighbor classifier

Q: With N examples,
how fast are training
and prediction?

A: Train $O(1)$,
predict $O(N)$

```

import numpy as np

class NearestNeighbor:
    def __init__(self):
        pass

    def train(self, X, y):
        """ X is N x D where each row is an example. Y is 1-dimension of size N """
        # the nearest neighbor classifier simply remembers all the training data
        self.Xtr = X
        self.ytr = y

    def predict(self, X):
        """ X is N x D where each row is an example we wish to predict label for """
        num_test = X.shape[0]
        # lets make sure that the output type matches the input type
        Ypred = np.zeros(num_test, dtype = self.ytr.dtype)

        # loop over all test rows
        for i in xrange(num_test):
            # find the nearest training image to the i'th test image
            # using the L1 distance (sum of absolute value differences)
            distances = np.sum(np.abs(self.Xtr - X[i,:]), axis = 1)
            min_index = np.argmin(distances) # get the index with smallest distance
            Ypred[i] = self.ytr[min_index] # predict the label of the nearest example

        return Ypred


```

Nearest Neighbor classifier

Q: With N examples, how fast are training and prediction?

A: Train $O(1)$,
predict $O(N)$


This is bad: we want classifiers that are **fast** at prediction; **slow** for training is ok



Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data


BAD: $K = 1$ always works perfectly on training data



Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data



Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data



Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data



Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data



Setting Hyperparameters

Idea #1: Choose hyperparameters that work best on the data

BAD: $K = 1$ always works perfectly on training data



Idea #2: Split data into **train** and **test**, choose hyperparameters that work best on test data

BAD: No idea how algorithm will perform on new data



Idea #3: Split data into **train**, **val**, and **test**; choose hyperparameters on val and evaluate on test

Better!



Setting Hyperparameters

Setting Hyperparameters


Your Dataset

Idea #4: Cross-Validation: Split data into **folds**, try each fold as validation and average the results

fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test
fold 1	fold 2	fold 3	fold 4	fold 5	test

Useful for small datasets, but not used too frequently in deep learning

Setting Hyperparameters




Example of 5-fold cross-validation for the value of **k**.

Each point: single outcome.


The line goes through the mean, bars indicated standard deviation

(Seems that $k \approx 7$ works best for this data)

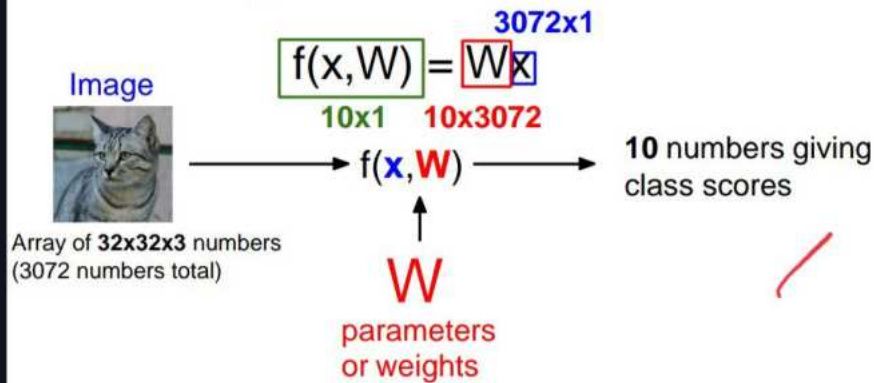
Parametric Approach



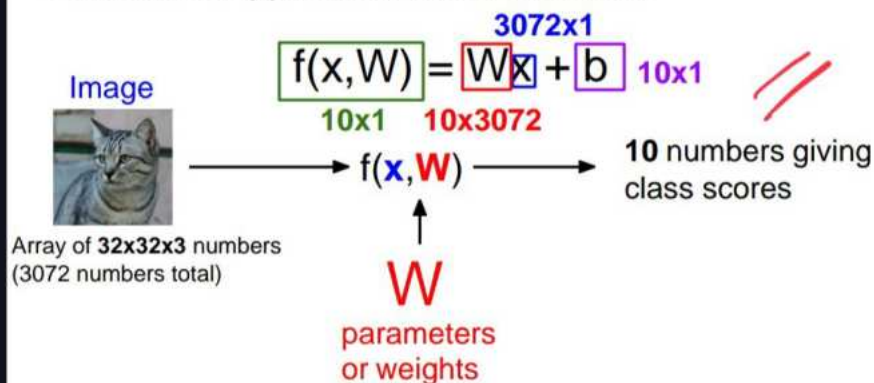
Parametric Approach: Linear Classifier




Parametric Approach: Linear Classifier



Parametric Approach: Linear Classifier



Example with an image with 4 pixels, and 3 classes (cat/dog/ship)



Interpreting a Linear Classifier



$$f(x, W) = Wx + b$$

What is this thing doing?




Interpreting a Linear Classifier




$$f(x, W) = Wx + b$$

Example trained weights
of a linear classifier
trained on CIFAR-10:



Interpreting a Linear Classifier



$$f(x, W) = Wx + b$$




Array of $32 \times 32 \times 3$ numbers
(3072 numbers total)

*Hyperplane?
Normal to a hyperplane*

Hard cases for a linear classifier


Class 1:
pixels coord > 0 odd

Class 2:
pixels coord > 0 even




Class 1:
 $1 \leq L2 \text{ norm} \leq 2$

Class 2:
Everything else



Class 1:
Three modes

Class 2:
Everything else



Convolution, feature map, pooling,
FCN

Basics of CNN (as taught in
class)

Thank you



Here are concise notes on Image Enhancement in the Frequency Domain for a 6-mark question:

1. Fourier Transform:

- Transforms a spatial image into its frequency domain representation.
- Used to process and enhance images by modifying their frequency components.

2. Filters in Frequency Domain:

- Low-Pass Filters (LPF): Smoothens images by removing high-frequency components.
 - Examples: Ideal LPF, Butterworth LPF, Gaussian LPF.
- High-Pass Filters (HPF): Sharpens images by removing low-frequency components.
 - Examples: Ideal HPF, Butterworth HPF, Gaussian HPF.

3. Notch Filters:

- Used to remove periodic noise by eliminating specific frequencies in the transform.

4. Selective Filters:

- Bandreject Filters: Removes a range of frequencies.
- Bandpass Filters: Keeps only a specific frequency band.

5. Image Sharpening:

- Achieved by subtracting the low-frequency components from the original image.

6. Applications:

- Used for noise removal, image smoothing, and sharpening.

Let me know if you'd like detailed explanations for any of these points.