

# Chapter 6: CPU Scheduling

---





# Chapter 6: CPU Scheduling

---

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Thread Scheduling
- Multiple-Processor Scheduling
- Real-Time CPU Scheduling
- Operating Systems Examples
- Algorithm Evaluation





# Objectives

---

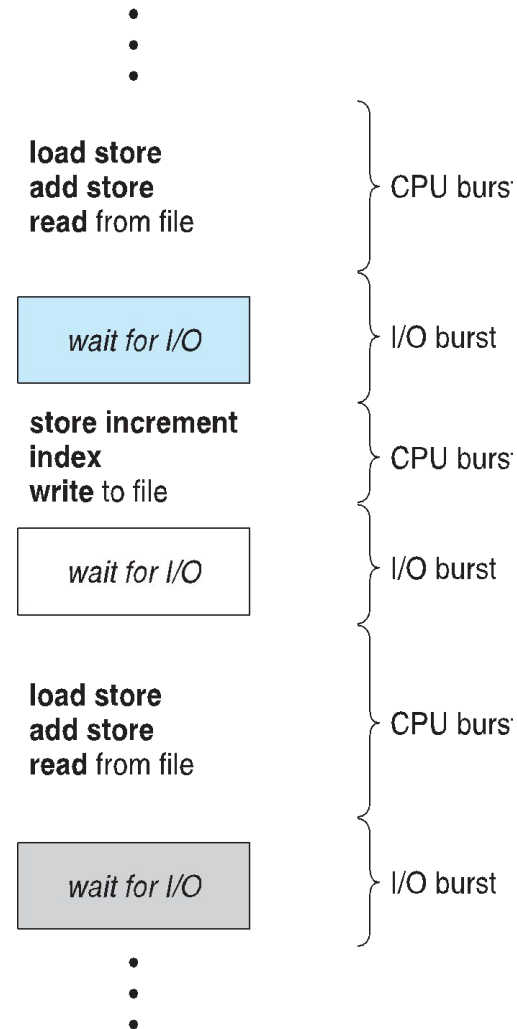
- To introduce CPU scheduling, which is the basis for multiprogrammed operating systems
- To describe various CPU-scheduling algorithms
- To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system
- To examine the scheduling algorithms of several operating systems





# Basic Concepts

- Maximum CPU utilization obtained with multiprogramming
- **CPU-I/O Burst Cycle** – Process execution consists of a **cycle** of **CPU execution** and **I/O wait**. Processes alternate between these two states.
- **CPU burst** followed by **I/O burst**
- CPU burst distribution is of main concern





# CPU Scheduler

- **Short-term scheduler** selects from among the processes in ready queue, and allocates the CPU to one of them
  - Queue may be ordered in various ways
- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state
  2. Switches from running to ready state
  3. Switches from waiting to ready
  4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**
  - Consider access to shared data
  - Consider preemption while in kernel mode
  - Consider interrupts occurring during crucial OS activities





# Dispatcher

- Dispatcher module **gives control of the CPU** to the process selected by the short-term scheduler; this involves:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program
- **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running





# Scheduling Criteria

---

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – Number of processes that complete their execution per time unit
- **Turnaround time** – amount of time needed to execute a particular process. It is the sum of the periods spent waiting to get into memory, waiting in the ready queue, executing on the CPU, and doing I/O.
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)





# Scheduling Algorithm Optimization Criteria

---

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time







# First- Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
----------------	-------------------

$P_1$	24
-------	----

$P_2$	3
-------	---

$P_3$	3
-------	---

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$  that arrive at time 0.

The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$





# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- **Convoy effect** - short process behind long process
  - Consider one CPU-bound and many I/O-bound processes





# Shortest-Job-First (SJF) Scheduling

---

- Associate with each process the length of its next CPU burst
  - Use these lengths to schedule the process with the shortest time
- SJF is optimal – gives minimum average waiting time for a given set of processes
  - The difficulty is knowing the length of the next CPU request
  - Could ask the user






# Example of SJF

	<u>Process</u>	<u>Burst Time</u>
$P_1$	6	
$P_2$	8	
$P_3$	7	
$P_4$	3	

- SJF scheduling chart



- Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$






# Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

	<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8	
$P_2$	1	4	
$P_3$	2	9	
$P_4$	3	5	

- Preemptive* SJF Gantt Chart



- Average waiting time =  $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$  msec





# Priority Scheduling

---

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - Preemptive
  - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Problem  $\equiv$  **Starvation** – low priority processes may never execute
- Solution  $\equiv$  **Aging** – as time progresses increase the priority of the process






# Example of Priority Scheduling

	<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3	
$P_2$	1	1	
$P_3$	2	4	
$P_4$	1	5	
$P_5$	5	2	

- Priority scheduling Gantt Chart



- Average waiting time = 8.2 msec





# Round Robin (RR)

---

- Each process gets a small unit of CPU time (**time quantum**  $q$ ), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once. No process waits more than  $(n-1)q$  time units.
- Timer interrupts every quantum to schedule next process
- Performance
  - $q$  large  $\Rightarrow$  FIFO
  - $q$  small  $\Rightarrow q$  must be large with respect to context switch, otherwise overhead is too high







# Example of RR with Time Quantum = 4


Process    Burst Time

$P_1$     24

$P_2$     3

$P_3$     3

- The Gantt chart is:




- Typically, higher average turnaround than SJF, but better *response*
- $q$  should be large compared to context switch time
- $q$  usually 10ms to 100ms, context switch  $< 10$  usec






# Time Quantum and Context Switch Time





# Turnaround Time Varies With The Time Quantum



process	time
$P_1$	6
$P_2$	3
$P_3$	1
$P_4$	7

80% of CPU bursts  
should be shorter than  $q$





# Multiprocessor Scheduling

---

Multiprocessor would earlier mean a system with multiple units of physical single-core CPUs. But in present times, a multiprocessor system may also refer to any of the following system architectures:

## a) Homogeneous multiprocessing

- i. a multi-core CPU or a set of multicore CPUs
- ii. Multi-threaded cores
- iii. Non-Uniform Memory Architecture (NUMA) systems


## b) Heterogeneous multiprocessing.



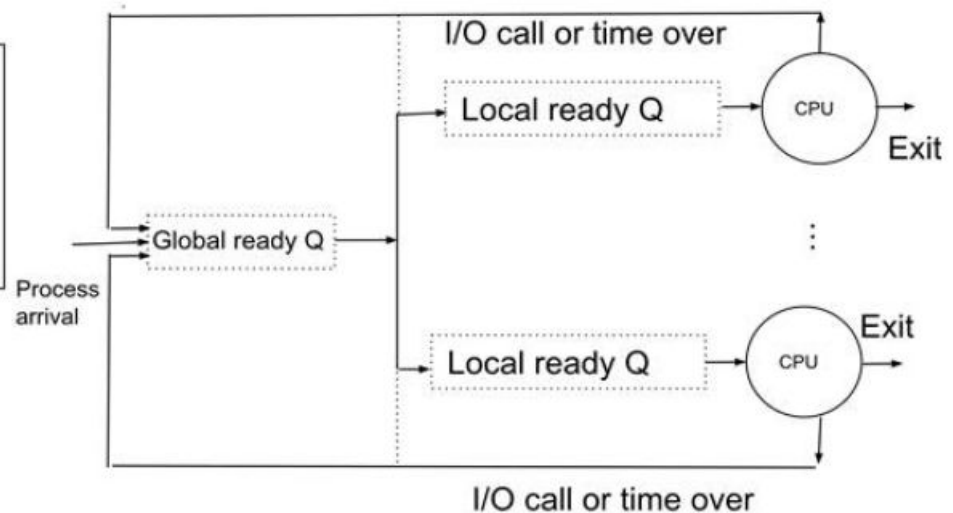


## i) Multiprocessor scheduling

- CPU scheduling more complex when multiple CPUs are available
- **Homogeneous processors** within a multiprocessor
  - **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing
  - **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes



a) Asymmetric multiprocessing



b) Symmetric multiprocessing





## i) *Multiprocessor scheduling*

---

Two important issues need to be discussed in relation to SMP systems:

- **Load balancing** attempts to keep workload evenly distributed
  - **Push migration** – periodic task checks load on each processor, and if found, pushes some tasks from overloaded CPU to other CPUs
  - **Pull migration** – idle processors pulls waiting task from busy processor
- **Processor affinity** – When the process (or thread) migrates to a different processor, the corresponding cache attached to the processor does not have the code, and data and it needs to store them again.
  - **soft affinity** - OS often attempts to schedule a given thread to a single processor, even though the allotment is not always guaranteed
  - **hard affinity** - allows processes to make system calls for scheduling to a given processor.





## ii) Multicore Processors


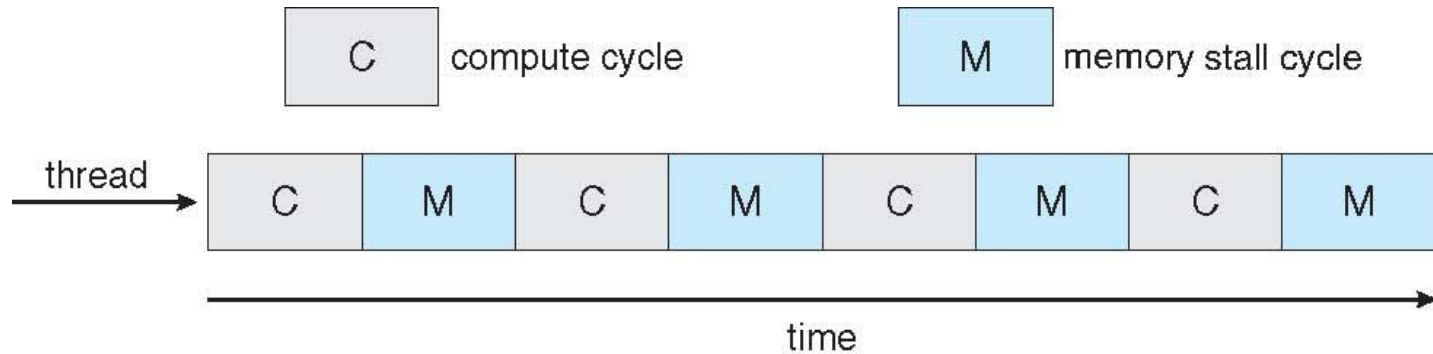
---

- Recent trend to place multiple processor cores on same physical chip
- Faster and consumes less power
- Multiple threads per core also growing
  - Takes advantage of **memory stall** to make progress on another thread while **memory retrieve** happens.
  - For a cache miss, a thread has to often wait for fetching data from memory (called **memory stall**) due to mismatch in speed of processing in the core and that of memory hardware. To utilize that idle time in the core, chip-level hardware threads are supported. When one hardware thread does **computation (C)**, another thread handles **memory stall (M)** in an interleaving fashion





# Multithreaded Multicore System








### *iii) Non-Uniform Memory Access(NUMA) and CPU Scheduling*

If OS scheduling and the memory management algorithms consider NUMA architecture, the threads can be allocated the CPU that is closest to the memory where the thread is loaded.



Note that memory-placement algorithms can also consider affinity





# Heterogeneous multiprocessing

---

- Some of the present day mobile devices use multicore processors of different processing attributes (clock speed, power requirement etc.). Such systems are called heterogeneous multiprocessing (HMP).
- This is mainly used to save battery power for long hours.
- For example, in ARM processors





# Exercise 1

- Consider the following set of processes, with the length of the CPU burst given in milliseconds:

Process	Burst	Time Priority
P1	2	2
P2	1	1
P3	8	4
P4	4	2
P5	5	3

- The processes are assumed to have arrived in the order P1, P2, P3, P4, P5, all at time 0.
- a. Draw four Gantt charts that illustrate the execution of these processes using the following scheduling algorithms: FCFS, SJF, nonpreemptive priority (a larger priority number implies a higher priority), and RR (quantum = 2).
- b. What is the turnaround time of each process for each of the scheduling algorithms in part a?
- c. What is the waiting time of each process for each of these scheduling algorithms?
- d. Which of the algorithms results in the minimum average waiting time (over all processes)?





## Exercise 2

- The following processes are being scheduled using a preemptive, round robin scheduling algorithm. Each process is assigned a numerical priority, with a higher number indicating a higher relative priority. In addition to the processes listed below, the system also has an *idle task* (which consumes no CPU resources and is identified as *Pidle*). This task has priority 0 and is scheduled whenever the system has no other available processes to run. The length of a time quantum is 10 units. If a process is preempted by a higher-priority process, the preempted process is placed at the end of the queue

Thread	Priority	Burst	Arrival
$P_1$	40	20	0
$P_2$	30	25	25
$P_3$	30	25	30
$P_4$	35	15	60
$P_5$	5	10	100
$P_6$	10	10	105

- Show the scheduling order of the processes using a Gantt chart.
- What is the turnaround time for each process?
- What is the waiting time for each process?
- What is the CPU utilization rate?



# End of Chapter 6

---

